

Image Processing Techniques Using OpenCV

Introduction to Image Transformations

Image transformations are fundamental operations in image processing that allow us to manipulate images for various purposes such as enhancing visual quality, extracting features, or preparing data for machine learning models. These transformations can be categorized into affine and non-affine transformations.

- **Affine Transformation:** A linear mapping method that preserves points, straight lines, and planes. It includes operations like translation, rotation, scaling, and shearing.

- **Non-Affine Transformation:** Non-linear transformations that distort the image in ways that do not preserve straight lines, such as perspective warping.

-Syntax

OpenCV provides several functions for applying transformations:

```
import cv2
```

```
import numpy as np
```

```
# Affine Transformation
```

```
M = np.float32([[1, 0, tx], [0, 1, ty]]) # Translation matrix
```

```
translated_image = cv2.warpAffine(image, M, (width, height))
```

```
# Rotation
```

```
M = cv2.getRotationMatrix2D(center, angle, scale)
```

```
rotated_image = cv2.warpAffine(image, M, (width, height))
```

Explanation:

- `cv2.warpAffine`: Applies an affine transformation to an image.

- `cv2.getRotationMatrix2D`: Computes the rotation matrix for a given center, angle, and scale.

Image Translations

Image translation involves moving an image up, down, left, or right without altering its size or orientation.

Translation shifts every pixel in the image by a specified amount along the x-axis (horizontal) and y-axis (vertical).

-Syntax

```
import cv2  
import numpy as np  
# Define translation matrix  
M = np.float32([[1, 0, tx], [0, 1, ty]]) # tx: horizontal shift, ty: vertical shift  
translated_image = cv2.warpAffine(image, M, (width, height))
```

-Example

```
image = cv2.imread('example.jpg')  
height, width = image.shape[:2]  
tx, ty = 50, 30 # Move 50 pixels right and 30 pixels down  
M = np.float32([[1, 0, tx], [0, 1, ty]])  
translated_image = cv2.warpAffine(image, M, (width, height))  
cv2.imshow('Translated Image', translated_image)  
cv2.waitKey(0)
```

-Explanation

The translation matrix `M` specifies how much to shift the image. Positive values of `tx` and `ty` move the image to the right and down, respectively.

Rotations and Horizontal Flipping

Rotating an image involves spinning it around a point, while flipping mirrors the image horizontally or vertically.

- Rotation: Spins the image around a specified center point by a given angle.
- Flipping: Reverses the image along the horizontal or vertical axis.

-Syntax

Rotation

```
M = cv2.getRotationMatrix2D(center, angle, scale)  
rotated_image = cv2.warpAffine(image, M, (width, height))
```

Flipping

```
flipped_image = cv2.flip(image, flip_code)
```

-Example

```
image = cv2.imread('example.jpg')  
height, width = image.shape[:2]  
  
# Rotate 90 degrees clockwise  
center = (width // 2, height // 2)  
M = cv2.getRotationMatrix2D(center, -90, 1)  
rotated_image = cv2.warpAffine(image, M, (width, height))  
  
# Horizontal Flip  
flipped_image = cv2.flip(image, 1)  
cv2.imshow('Rotated Image', rotated_image)  
cv2.imshow('Flipped Image', flipped_image)  
cv2.waitKey(0)
```

Scaling, Resizing, and Interpolation

Scaling changes the size of an image, while interpolation determines how missing pixel values are calculated during resizing.

- Scaling: Adjusts the dimensions of an image (e.g., enlarging or shrinking).
- Interpolation: Estimates new pixel values when an image is resized.

-Syntax

```
resized_image = cv2.resize(image, (new_width, new_height),  
interpolation=cv2.INTER_LINEAR)
```

-Example

```
image = cv2.imread('example.jpg')
```

```
# Resize to half the original size
```

```
resized_image = cv2.resize(image, None, fx=0.5, fy=0.5,  
interpolation=cv2.INTER_AREA)
```

```
cv2.imshow('Resized Image', resized_image)
```

```
cv2.waitKey(0)
```

-Explanation

- `cv2.INTER_LINEAR`: Suitable for shrinking images.
- `cv2.INTER_CUBIC`: Better for enlarging but slower.
- `cv2.INTER_AREA`: Best for shrinking.

Image Pyramids

Image pyramids represent an image at multiple resolutions, useful for tasks like object detection.

Image pyramids reduce or expand an image's resolution iteratively.

-Syntax

Gaussian Pyramid

```
lower_res = cv2.pyrDown(image)  
higher_res = cv2.pyrUp(lower_res)
```

-Example

```
image = cv2.imread('example.jpg')
```

Downscale using Gaussian Pyramid

```
lower_res = cv2.pyrDown(image)
```

Upscale back

```
higher_res = cv2.pyrUp(lower_res)
```

```
cv2.imshow('Lower Resolution', lower_res)  
cv2.imshow('Higher Resolution', higher_res)  
cv2.waitKey(0)
```

Cropping

Cropping extracts a specific region of interest (ROI) from an image.

Cropping selects a rectangular portion of an image by specifying pixel ranges.

-Syntax

```
cropped_image = image[y1:y2, x1:x2]
```

-Example

```
image = cv2.imread('example.jpg')
```

Crop a region

```
cropped_image = image[50:200, 100:300] # y1:y2, x1:x2
```

```
cv2.imshow('Cropped Image', cropped_image)
cv2.waitKey(0)
```

Arithmetic and Bitwise Operations

Arithmetic and bitwise operations manipulate pixel values to adjust brightness or create masks.

-Arithmetic Operations: Add/subtract scalar values to brighten/darken images.

- Bitwise Operations: Perform logical operations on binary masks.

-Syntax

Brightening/Darkening

```
brightened_image = cv2.add(image, scalar_value)
```

```
darkened_image = cv2.subtract(image, scalar_value)
```

Bitwise AND

```
result = cv2.bitwise_and(image1, image2, mask=mask)
```

-Example

```
image = cv2.imread('example.jpg')
```

```
# Brighten by adding 50 to each pixel
```

```
brightened_image = cv2.add(image, 50)
```

```
# Create a mask and apply bitwise AND
```

```
mask = np.zeros(image.shape[:2], dtype=np.uint8)
```

```
mask[50:200, 100:300] = 255
```

```
masked_image = cv2.bitwise_and(image, image, mask=mask)
```

```
cv2.imshow('Brightened Image', brightened_image)
```

```
cv2.imshow('Masked Image', masked_image)
```

```
cv2.waitKey(0)
```

```
...
```

Blurring - The Many Ways We Can Blur Images & Why It's Important

Blurring reduces image noise and detail by averaging pixel values in a neighborhood. It is crucial for preprocessing tasks like edge detection, noise removal, and reducing high-frequency details.

-Types of Blurring Techniques

1. Averaging Blur: Uses a normalized box filter to average pixel values.
2. Gaussian Blur: Applies a weighted average with a Gaussian kernel.
3. Median Blur: Replaces each pixel with the median value of its neighbors.
4. Bilateral Filtering: Preserves edges while blurring.

-Syntax

Averaging Blur

```
blurred_image = cv2.blur(image, (kernel_size, kernel_size))
```

Gaussian Blur

```
blurred_image = cv2.GaussianBlur(image, (kernel_size, kernel_size), sigmaX)
```

Median Blur

```
blurred_image = cv2.medianBlur(image, kernel_size)
```

Bilateral Filtering

```
blurred_image = cv2.bilateralFilter(image, d, sigmaColor, sigmaSpace)
```

-Example

```
import cv2
```

```
image = cv2.imread('example.jpg')
```

Averaging Blur

```
averaged_blur = cv2.blur(image, (5, 5))
```

Gaussian Blur

```
gaussian_blur = cv2.GaussianBlur(image, (5, 5), 0)
```

Median Blur

```
median_blur = cv2.medianBlur(image, 5)
```

Bilateral Filtering

```
bilateral_blur = cv2.bilateralFilter(image, 9, 75, 75)
```

```
cv2.imshow('Averaged Blur', averaged_blur)
```

```
cv2.imshow('Gaussian Blur', gaussian_blur)
```

```
cv2.imshow('Median Blur', median_blur)
```

```
cv2.imshow('Bilateral Blur', bilateral_blur)
```

```
cv2.waitKey(0)
```

-Explanation

-`cv2.blur`: Simple averaging of pixel values.

-`cv2.GaussianBlur`: Weighted average with a Gaussian kernel, ideal for smoothing.

-`cv2.medianBlur`: Removes salt-and-pepper noise effectively.

-`cv2.bilateralFilter`: Smooths while preserving edges, useful for beautification.

Sharpening - Reverse Your Image Blurs

Sharpening enhances edges and fine details in an image, often reversing the effects of blurring. It increases contrast between adjacent pixels.

-Technique

Sharpening is achieved using a **kernel-based convolution** operation. A common sharpening kernel is:

[0, -1, 0]

[-1, 5, -1]

[0, -1, 0]

-Syntax

```
sharpened_image = cv2.filter2D(image, -1, kernel)
```

-Example

```
import cv2
```

```
import numpy as np
```

```
image = cv2.imread('example.jpg')
```

```
# Define a sharpening kernel
```

```
kernel = np.array([[0, -1, 0],
```

```
                  [-1, 5, -1],
```

```
                  [0, -1, 0]])
```

```
# Apply sharpening
```

```
sharpened_image = cv2.filter2D(image, -1, kernel)
```

```
cv2.imshow('Original Image', image)
```

```
cv2.imshow('Sharpened Image', sharpened_image)
```

```
cv2.waitKey(0)
```

-Explanation

- `cv2.filter2D`: Applies a custom convolution kernel to the image.

- The kernel emphasizes edges by increasing pixel intensity differences.

Thresholding (Binarization) - Making Certain Image Areas Black or White

Thresholding converts a grayscale image into a binary image by setting pixel values above or below a threshold to either black (0) or white (255). It is widely used in segmentation and object detection.

-Types of Thresholding

1. Simple Thresholding: Fixed threshold value.
2. Adaptive Thresholding: Dynamically calculates thresholds based on local regions.
3. Otsu's Thresholding: Automatically determines the optimal threshold for bimodal images.

-Syntax

Simple Thresholding

```
_, binary_image = cv2.threshold(gray_image, threshold_value, max_value,  
threshold_type)
```

Adaptive Thresholding

```
binary_image = cv2.adaptiveThreshold(gray_image, max_value,  
adaptive_method, threshold_type, block_size, C)
```

Otsu's Thresholding

```
_, binary_image = cv2.threshold(gray_image, 0, max_value,  
cv2.THRESH_BINARY + cv2.THRESH_OTSU)
```

-Example

```
import cv2
```

```
image = cv2.imread('example.jpg', cv2.IMREAD_GRAYSCALE)
```

Simple Thresholding

```
_, simple_thresh = cv2.threshold(image, 127, 255, cv2.THRESH_BINARY)
```

Adaptive Thresholding

```
adaptive_thresh = cv2.adaptiveThreshold(image, 255,  
cv2.ADAPTIVE_THRESH_MEAN_C, cv2.THRESH_BINARY, 11, 2)
```

```
cv2.imshow('Simple Thresholding', simple_thresh)
```

```
cv2.imshow('Adaptive Thresholding', adaptive_thresh)
```

```
cv2.waitKey(0)
```

-Explanation

-Simple Thresholding: Converts pixels above a fixed threshold to white and others to black.

- Adaptive Thresholding: Useful for uneven lighting conditions.

Github Repo with all image processing code examples:

**[https://github.com/rajm-webocult/Basics-of-OpenCV/tree/main/
image_processing_code](https://github.com/rajm-webocult/Basics-of-OpenCV/tree/main/image_processing_code)**